DevOps for Cloud Assignment (S2-24_CCZG507)

Student Details:

- BITS ID: 2024MT03053
- Student Name: Balasubramaniyan Murugappa

This project demonstrates a full CI/CD and monitoring setup for a backend application using FastAPI, Docker, Kubernetes, and Prometheus.



Task 1: Create Backend using FastAPI

- Endpoint: / or /get_info
- Returns environment variables APP_VERSION and APP_TITLE along with the pod name (load balancer)
- Requirements

```
fastapi
uvicorn
python-dotenv
prometheus_client
psutil
```

```
from fastapi import FastAPI, Response
from fastapi.responses import RedirectResponse
```

```
from prometheus_client import (
    Counter,
    generate_latest,
    Gauge,
    CONTENT TYPE LATEST,
)
from dotenv import load_dotenv
import os
import psutil
import threading
import time
import random
import socket
load dotenv()
app = FastAPI()
REQUEST COUNT = Counter("get info requests total", "Total number of
GET /get info requests")
REQUEST_COUNT_PER_VERSION = Counter(
    "get_info_requests_total_by_version",
    "GET /get_info requests by app version",
    ["version"]
CPU_USAGE = Gauge("cpu_usage_percent", "CPU usage percentage")
MEMORY_USAGE = Gauge("memory_usage_percent", "Memory usage
percentage")
UPTIME = Gauge("uptime_seconds", "App uptime in seconds")
THREAD_COUNT = Gauge("thread_count", "Number of active threads")
DISK_USAGE = Gauge("disk_usage_percent", "Disk usage percentage")
START_TIME = time.time()
@app.get("/")
def root():
    return RedirectResponse(url="/get_info")
@app.get("/get_info")
async def get_info():
    app_title = os.getenv("APP_TITLE", "My FastAPI App")
    app_version = os.getenv("APP_VERSION", "1.0")
    REQUEST_COUNT.inc()
    REQUEST_COUNT_PER_VERSION.labels(version=app_version).inc()
    return {
        "APP_VERSION": app_version,
        "APP_TITLE": app_title,
        "MESSAGE": "Hello from " + socket.gethostname(),
    }
@app.get("/metrics")
def metrics():
    CPU_USAGE.set(psutil.cpu_percent())
```

```
MEMORY_USAGE.set(psutil.virtual_memory().percent)
UPTIME.set(time.time() - START_TIME)
THREAD_COUNT.set(threading.active_count())
DISK_USAGE.set(psutil.disk_usage("/").percent)
return Response(generate_latest(), media_type=CONTENT_TYPE_LATEST)
```

```
venv ~/workspace/Mtech/2024mt03053_dc (20.314s)
uvlcorn app.mmain:app --host 127.0.0.1 --port 8000 --reload

INFO: Will watch for changes in these directories: ['/Users/dgtalbug/workspace/Mtech/2024mt03053_dc']
INFO: Uvlcorn running on http://127.0.0.1:8000 (Press CTRL+C to quit)
INFO: Started reloader process [7457] using StatReload
INFO: Started server process [7457]
INFO: Waiting for application startup.
INFO: Application startup complete.
INFO: 127.0.0.1:58514 - "GET / HTTP/1.1" 307 Temporary Redirect
INFO: 127.0.0.1:58514 - "GET / get_info HTTP/1.1" 200 0K
```

```
● ● ② 127.0.0.1:52731/get_info × +

← → ♂ ⑥ 127.0.0.1:52731/get_info 

Pretty print ☑

{
"APP_VERSION": "1.0",
"APP_TITLE": "2024MT9353-DevOps-Assignment",
"MESSAGE": "Hello from fastapi-deployment-77d947d545-dxt5d"
}
```

Task 2: Dockerize the Application

• Dockerfile created to containerize the FastAPI app

```
FROM python:3.11-slim

WORKDIR /app

COPY app/requirements.txt .

RUN pip install --no-cache-dir -r requirements.txt

COPY . .

EXPOSE 8000

CMD ["uvicorn", "app.main:app", "--host", "0.0.0.0", "--port", "8000", "--workers", "2"]
```

• Image built using: docker build -t img-2024mt03053 .

```
| venv -/workspace/Mtech/2024mt03053_dc (27.572s) | docker butld -t ing-2024mt03053_dc (27.572s) | docker butld -t ing-
```

Task 3: Run Docker Container

• Run using:

```
docker run -d --name CNR-2024MT03053 -p 8000:8000 -e APP_VERSION=1.0 - e APP_TITLE="2024MT0353-DevOps-Assignment" img-2024mt03053
```

Task 4: Deploy Image to Kubernetes (Minikube)

- Created ConfigMap with APP_VERSION and APP_TITLE
- Deployment created with 2 replicas using the image img-2024mt03053
- · Service exposed with loadbalancer

```
apiVersion: v1
kind: ConfigMap
metadata:
   name: config-2024mt03053
data:
   APP_VERSION: '1.0'
   APP_TITLE: '2024MT0353-Dev0ps-Assignment'
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
   name: fastapi-deployment
spec:
   replicas: 3
   selector:
     matchLabels:
        app: fastapi-app
template:
     metadata:
```

```
labels:
    app: fastapi-app
spec:
  containers:
    - name: fastapi-container
      image: img-2024mt03053:latest
      imagePullPolicy: Never
      ports:
        - containerPort: 8000
      envFrom:
        - configMapRef:
            name: config-2024mt03053
      readinessProbe:
        httpGet:
          path: /
          port: 8000
        initialDelaySeconds: 5
        periodSeconds: 10
        failureThreshold: 3
        successThreshold: 1
```

```
apiVersion: v1
kind: Service
metadata:
   name: fastapi-service
spec:
   type: LoadBalancer
   selector:
     app: fastapi-app
   sessionAffinity: None
   ports:
     - protocol: TCP
     port: 80
     targetPort: 8000
```

• Deploy & check the kube status

```
kubectl apply -f k8s/ConfigMap.yaml
kubectl apply -f k8s/Deployment.yaml
kubectl apply -f k8s/Service.yaml
minikube dashboard
```

Deployment

```
Venv ~/workspace/Mtech/2024mt03053 dc @ minikube (0.828s)
kubectl apply -f k8s/ConfigMap.yaml -n app && kubectl apply -f k8s/Deployment.yaml -n app && kubectl apply -f k8s/Service.yaml -n app
configmap/config-2024mt03053 created
deployment.apps/fastapi-deployment created
service/fastapi-service created

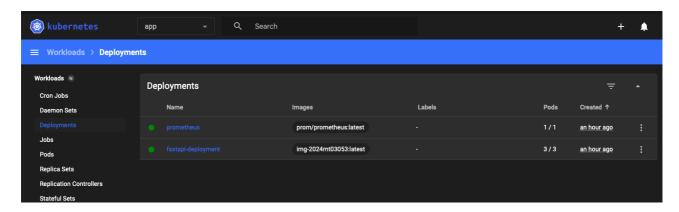
venv ~/workspace/Mtech/2024mt03053 dc @ minikube (0.839s)
kubectl apply -f prometheus/ConfigMap.yaml -n app && kubectl apply -f prometheus/Deployment.yaml -n app && kubectl apply -f prometheus/Service.yaml -n app
configmap/prometheus-config created
deployment.apps/prometheus created
service/prometheus restrice created

venv ~/workspace/Mtech/2024mt03053.dc (45m 50.62s)
minikube service fastapi-service --url -n app
http://127.0.0.1:52157

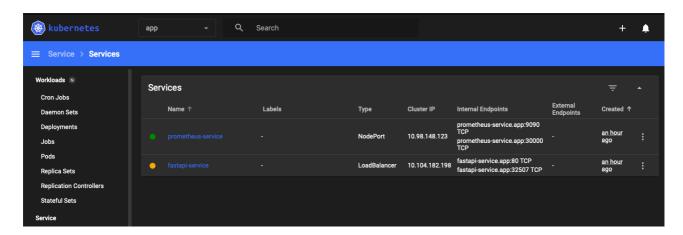
1 Because you are using a Docker driver on darwin, the terminal needs to be open to run it.
```

Status

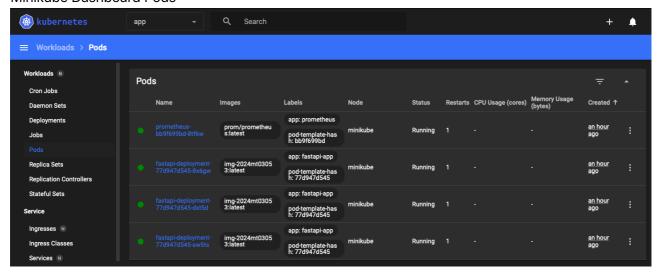
• Minikube Dashboard Deployments



• Minikube Dashboard Service



Minikube Dashboard Pods



Task 5: Configure LoadBalancer

- Exposed app using a LoadBalancer service
- Verified service endpoint using:

```
minikube service fastapi—service ——url
```

Task 6: Configure Prometheus

- V Installed Prometheus and configured it with Kubernetes using a ConfigMap.
- V Enabled system and application-level metrics via prometheus_client.
- **V** Exposed a /metrics endpoint on the FastAPI app for Prometheus scraping.

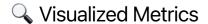
```
apiVersion: v1
kind: ConfigMap
metadata:
   name: prometheus-config
data:
   prometheus.yml: |
     global:
     scrape_interval: 15s
   scrape_configs:
     - job_name: 'fastapi'
     static_configs:
     - targets: ['fastapi-service:80']
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: prometheus
spec:
  replicas: 1
  selector:
    matchLabels:
      app: prometheus
  template:
    metadata:
      labels:
        app: prometheus
    spec:
      containers:
        - name: prometheus
          image: prom/prometheus:latest
          ports:
            - containerPort: 9090
          volumeMounts:
            - name: config-volume
              mountPath: /etc/prometheus/
      volumes:
        - name: config-volume
          configMap:
            name: prometheus-config
```

```
apiVersion: v1
kind: Service
metadata:
   name: prometheus-service
spec:
   type: NodePort
   selector:
    app: prometheus
ports:
   - port: 9090
     targetPort: 9090
     nodePort: 30000
```

• Deploy & check the kube status

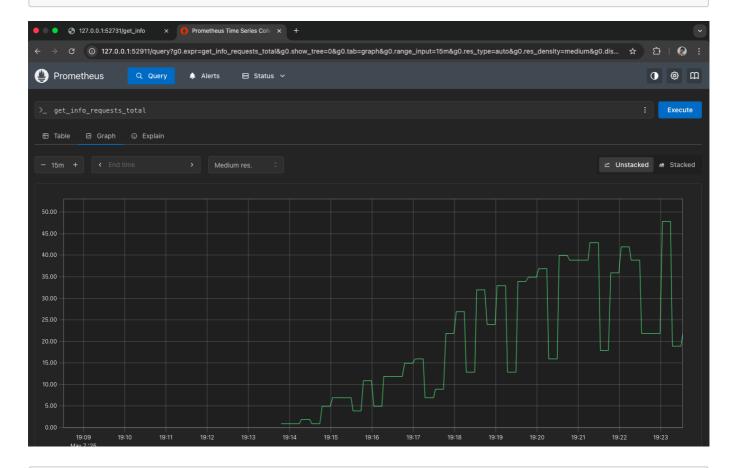
```
kubectl apply -f prometheus/ConfigMap.yaml
kubectl apply -f prometheus/Deployment.yaml
kubectl apply -f prometheus/Service.yaml
```



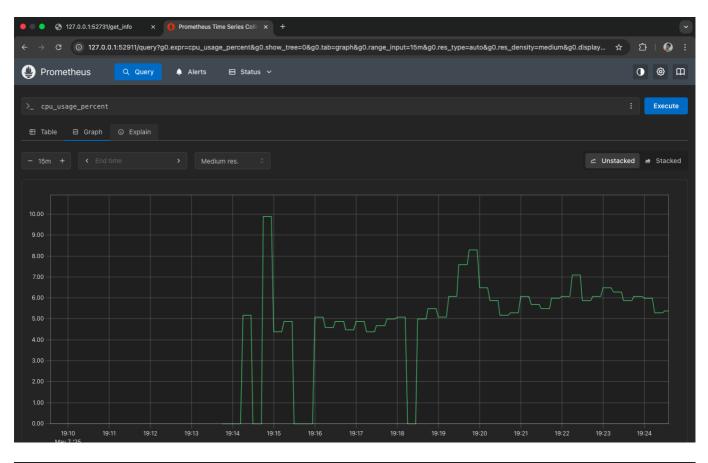
Metric Name	Description
get_info_requests_total	Total number of GET /get_info requests
<pre>get_info_requests_total_by_version</pre>	Requests to /get_info labeled by app version
cpu_usage_percent	Current CPU usage percentage
memory_usage_percent	Current memory usage percentage
uptime_seconds	Uptime of the application in seconds
thread_count	Number of active threads
disk_usage_percent	Disk usage percentage of root filesystem

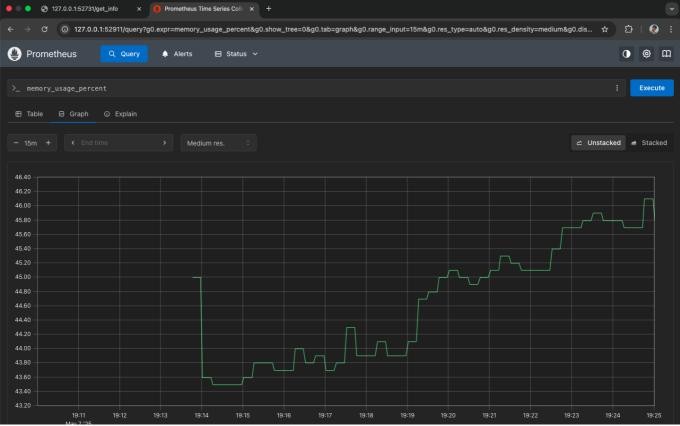
Sample Prometheus Queries

Total requests
get_info_requests_total

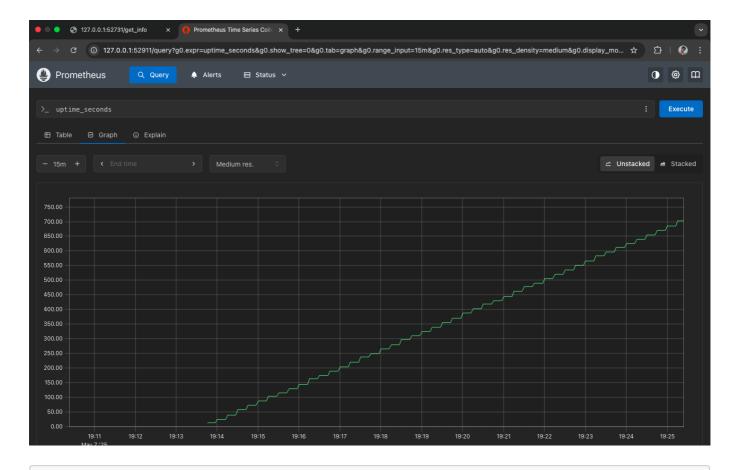


Current CPU and memory usage
cpu_usage_percent
memory_usage_percent

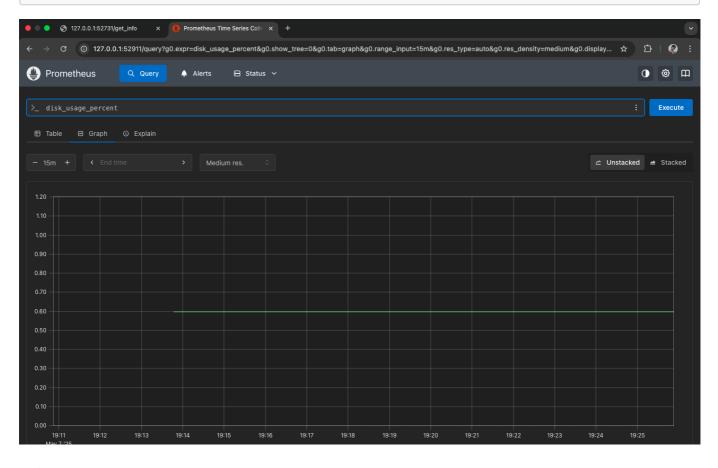




App uptime
uptime_seconds



Disk usage
disk_usage_percent



Accessing Prometheus

If running via NodePort:

```
http://<your-node-ip>:30000
```

Or using kubectl port-forward:

```
kubectl port-forward svc/prometheus-service 9090:9090
```

Then open your browser to:

```
http://localhost:9090
```

Directory Structure

```
ROLL_NUMBER_dc/

    Dockerfile

  README.md
  — арр
     — main.py
     — requirements.txt
 – k8s
    ConfigMap.yaml
      Deployment.yaml
    └─ Service.yaml
  - prometheus
    ─ ConfigMap.yaml
      Deployment.yaml
    Service.yaml
  - screenshots
     — app-response.png
      - application-init.png
      application-running-local.png
      — cpu−usage.png
      - curl-logs.png
      - curl-test.png
      — disk−usage.png
     — docker-image-build.png
      get-total-requests.png
      kube-deployments.png
      – kube-pods.png
      - kube-services.png
      kubectl-pods-deploy.png
      kubectl-pods-status.png
      - memory-usage.png
```

minikube-start.png
 target-status.png
 uptime-seconds.png

Challenges faced

- Helm installation: Tried Helm but faced some env issues
- K3s: Lightweight K8s again env issues
- Minikube: worked in both linux and mac.